

# UMANG SHAH

Data Engineer – New York

umangshah97@gmail.com | umangshah.me | linkedin.com/in/shah-umang-me | github.com/umangshah-js

## Summary

Data Engineer with 4+ years building scalable ETL, real-time streaming, and cloud-native platforms across fintech and healthcare.

## Education

**New York University** – M.S. in Computer Science, GPA: 3.96/4.0

Sep 2022 – May 2024

**VIT Chennai** – B.Tech in Computer Science and Engineering, GPA: 9.16/10

Jul 2015 – May 2019

## Experience

**Analyst (Data / Software Engineering)**, Nomura Holdings America – New York, NY

Jul 2024 – Present

- Delivered 67+ end-to-end features on a 15+ year .NET C# pricing platform by modernizing payoff logic and trade-pricing workflows.
- Added buffer-coupon pricing to provide fast provisional charge estimates on incoming trades, enabling desks to act on quotes before final accurate valuation completed.
- Reduced manual trader workflows by replacing Excel-based pricing with request-driven pipelines using hybrid REST and pub-sub compute.
- Automated the 144A post-trade lifecycle with bulk import, instrument creation, booking-line management, and contract generation.
- Improved pricing accuracy and trade visibility by integrating Euclid APIs, Radar, and Bloomberg into front-to-back workflows.

**Software Engineer Intern**, MIO Partners (McKinsey Investment Office) – New York, NY

Jun 2023 – Aug 2023

- Improved ETL reliability and developer velocity by migrating ActiveBatch/Talend workflows to AWS Step Functions + S3 and building a portal for logs, trigger history, and one-click runs.

**Senior Software Engineer (Data Engineering)**, UnitedHealth Group – Hyderabad, India

Jul 2019 – Jul 2022

- Standardized healthcare onboarding by building DXP pipelines ingesting flat files, CCD XML, HL7 messages, and Kafka streams into partitioned Hive Parquet datasets.
- Improved data quality and reliability through ingestion validation, Airflow orchestration, and Python/Angular controls for client scheduling and source formats.
- Delivered MTD/YTD analytics pipelines across 26M+ members on large Hive/Spark clusters; used single-query Hive for straightforward workloads and spearheaded PySpark jobs for complex transformations, tuning memory by queue and node capacity for stable throughput.
- Reduced near-real-time data SLA from 6 hours to 10 minutes by designing and deploying Kafka-based streaming pipelines.
- Reduced infrastructure cost and production downtime by migrating core services to Kubernetes.
- Improved incident detection and recovery by implementing Grafana and Prometheus monitoring/alerting for mission-critical workflows.

## Skills

**Data Engineering:** Spark, Kafka, Hive, HDFS, Snowflake, ETL, PostgreSQL, SQL, MongoDB, HBase, Redis, KDB+

**Cloud & DevOps:** AWS, Azure, Kubernetes, Docker, CI/CD, Jenkins, Nginx

**Web & Backend:** Node.js, Flask, .NET Framework, FastAPI

**Analytics & ML:** Pandas, Tableau, D3.js, Chart.js, NLTK, StanfordNLP

**Programming & Version Control:** Python, JavaScript, Java, C#, Git

## Projects

**Media Magic (Live Demo)**

Sep 2023 – Dec 2023

- Built a developer-first image delivery platform on AWS (S3 → SQS → Lambda → DynamoDB) plus SDK/CLI tooling, enabling path-based integration and sub-200ms image responses.
- Improved visual consistency and processing efficiency by using OpenCV + YOLO smart-focus cropping, then adding hash-based idempotency and lease-based request coalescing to prevent duplicate transforms.
- Improved delivery reliability and observability by orchestrating CloudFront-backed on-demand WebP transforms with readiness-aware retries and analytics for throughput, transform latency, and cache-hit efficiency.

**Forest Fire Tracker (GitHub)**

Feb 2023 – Apr 2023

- Built a Kafka-driven streaming pipeline to process Unity-generated satellite-image grids across ingestion, preprocessing, and analytics stages.
- Generated near-real-time fire-risk progression states by implementing OpenCV-based tree/cluster detection, fire-region extraction, and iterative spread prediction.
- Delivered reliable, live operational visibility by orchestrating Redis state management, Dask aggregation, and a Dash metrics dashboard with flow-control signaling.

## Awards

- Won “Open Gov Data Hack” 2018, a national-level competition organized by NIC, Government of India.
- Won best project award in Technical Expo’18, VIT Chennai.